

# TOM MARSH

Boston, MA | [LinkedIn](#) | [GitHub](#) | [Personal Website](#)

**SUMMARY** Senior Machine Learning Engineer with 8 years building production ML systems at TB/PB scale. Track record delivering measurable business impact: foundational ML infrastructure generating €15M+ annually, platform migrations saving \$200K+ yearly, and 0-1 startup ML system contributing to \$10M+ valuation. Physics background (MSc) with end-to-end ownership from architecture to deployment.

---

## EXPERIENCE

### Ströer Labs NZ - Senior Machine Learning Engineer

*Christchurch, New Zealand*

**Feb 2018 - Nov 2024**

- Reduced ad impression render time by 50% through ML inference optimisation, generating €15M+ additional annual revenue across 3B+ daily ad requests (1T+ annually)
- Led complete migration of ML data pipelines from Spark to DBT, eliminating Spark cluster costs (\$10K+/month), reducing pipeline runtime 85% (20 min to 3 min), and simplifying maintenance for successor team
- Architected migration from AWS SageMaker to custom Kubernetes infrastructure, enabling GPU training for LightGBM models (previously impossible), unifying ML platform with company infrastructure
- Cut cloud egress costs by 35% (\$200K+ annually) for product vertical through ML optimisation; maintenance team expanded savings to \$30K/month post-transition
- Led 0-1 development of multiple ML products (classification, time-series forecasting, optimization) from concept to production on PB-scale datasets, establishing technical architecture before transitioning to dedicated teams
- Drove ML platform modernization during company's bare-metal-to-AWS migration, collaborating with DevOps and Data Engineering to redesigndatalakes, account structures, and deployment practices
- Designed technical interview process for ML Engineering roles: created programming assessments, architecture evaluation framework, and systems design questions used across MLE/SWE/DevOps hiring
- Conducted 80+ technical interviews over 6 years, contributing to 15 hires as NZ office grew from 20 to 30 engineers
- Hosted interns each summer and mentored juniors through intermediate

### Pyper Vision - Founding Machine Learning Engineer (Contracts, part-time)

*Contract 1 | Christchurch, New Zealand*

**Jul 2024 - Feb 2025**

- Established complete ML infrastructure as sole technical hire, enabling aerospace startup pivot from hardware to ML-based fog forecasting (company expanded 5 to 20+ airports, valued \$10M+ post-funding)
- Architected end-to-end system on GCP from scratch: automated data collection, preprocessing pipelines, model training infrastructure, and production inference API with <100ms p99 latency
- Evaluated and deployed model architectures (LightGBM GBDT, TensorFlow CNNs) achieving commercially competitive forecast accuracy, delivering production system in 8 months
- System scaled to UK market expansion (advised on integration strategy post-contract)

- Designed physics-validated fog classification system using pseudo-labeling pipeline derived from meteorological research papers
  - Implemented clustering analysis by airport fog patterns and model error modes, creating systematic debugging framework for operations team to diagnose regional forecast failures
- 

## EDUCATION

**MSc, Physics - University of Canterbury** **Jul 2022 - Jul 2024**

- Thesis: Extreme precipitation events under climate change using 4000+ years of simulation data at multiple resolutions
- Collaboration with National Institute of Water and Atmospheric Research (NIWA)
- Presented research at Meteorological Society of New Zealand Annual Conference (2023)

**BSc, Physics & Computer Science - University of Canterbury** **Feb 2015 - Nov 2017**

- Capstone: Y<sub>2</sub>SiO<sub>5</sub> thin film growth for quantum information devices using Pulsed Laser Deposition
- 

## TECHNICAL SKILLS

**ML/Data:** PyTorch, TensorFlow, LightGBM, XGBoost, Scikit-learn, Pandas, NumPy, Spark, MLFlow, DBT

**Infrastructure:** AWS (SageMaker, EMR, S3, Glue, Athena), GCP (Cloud Functions, GCS), Docker, Kubernetes

**Data Engineering:** Apache Airflow, Kafka, SQL (Postgres, Presto, Trino), time-series databases

**Languages:** Python, SQL, Java, Scala, Kotlin, Go

**Miscellaneous:** Unix, Source Control (Git)